

rhema.ai

Design & UX documentation

What we built, why every decision was made, which references and rules we followed, the guardrails that protect people, and every test we ran — in one place.

Product

Faith-vocabulary app + pastor care

Audience

Gen Z / Gen Alpha, South & Southeast Asia; pastors in sensitive regions

Platforms

Web (primary), phone-first

Status

UI/UX complete · backend contract ready · human tests pending

Contents

1. Product summary
2. Problem, audience, principles
3. The three layers
4. Screen inventory (what · why)
5. Design system
6. Motion & emotional design
7. Onboarding, guest-first, landing, tour
8. Guardrails & safety
9. Moderation pipeline
10. Privacy, alert mode, hidden features
11. Languages & low bandwidth
12. References & how we used them
13. Rules we followed (videos, transcripts, guides)
14. Testing: X-ray, audits, AI, traditional
15. Benchmarks & scores
16. Originality check
17. Beta, feedback, analytics
18. Backend contract & handoff
19. Human flourishing by design
20. Open build & license
21. Shared schema & MCP server
22. Published eval set & rubric
23. Reusable patterns
24. Practitioner evidence
25. What didn't work
26. What's left
27. File map

1. Product summary

rhema.ai is a calm, private faith-vocabulary app. A reader looks up one word — *karma*, *salvation*, *moksha* — and sees how Hindu, Buddhist and Christian traditions each use it, side by side, never ranked, each described from its own texts with sources. A hidden **Faith mode** offers a deeper Christian bridge for those who look for it. Once a month, everyone may answer one anonymous question; answers become a published **Ideas map**. Behind a separate door, **pastors** in sensitive regions get light check-ins, a human mentor and human reviewers — the assistant drafts, a person always decides.

One-line pitch. One word. Three traditions. In your language — and a quiet place where pastors under pressure are held by people, not scored by software.

2. Problem, audience, principles

Problem

- Young people in multi-faith societies meet words like *karma* and *grace* constantly, but explanations online are either one-sided, preachy or academic.
- Pastors in sensitive regions carry heavy loads alone; existing church software (built for US megachurches) tracks attendance and money, not wellbeing — and can put them at risk if seen.

Audience

Who	Needs	What it means for design
Gen Z / Gen Alpha readers (India, Bangladesh, Nepal, Myanmar, Cambodia)	Fast, fair, private answers; no sign-up wall; their language	Guest-first, 6 languages, plain words (grade $\leq$ 8), calm dark visual, no hype
Pastors & lay leaders	Support without surveillance; works on shared, low-signal phones	"When you can" check-ins, no streaks or scores, offline queue, quick-hide, crisis line, mentor
Reviewers & mentors	Clear decisions, no names	Code names, one review at a time, AI draft clearly labelled
Leaders / admins	Safety switch; publish control	Region alert with typed confirm; map changes only on publish
US churches (optional)	Use our features inside Planning Center	Optional connection; read-only embed card; no finance

Core principles

1. **Understand, don't convince.** Traditions are described in their own words. No winner, no ranking.
2. **People decide.** The assistant drafts; a named role (reviewer, mentor, admin) approves. A stage never moves by AI.
3. **Private by default.** No names on public screens, no feed, no comments or votes, no finance.
4. **Calm, not flashy.** Earthy colour, slow motion, one purpose per screen.
5. **Fail in the open.** Every state is designed: loading, empty, error, offline, unavailable, blocked.
6. **Safe where it matters.** Sensitive features can be hidden or paused per region.

3. The three layers

Layer 1 — Dictionary (public)

- **Search** → **Word page**. Headword, one or two plain sentences, tradition labels, sources and a trust line ("3 sources · checked by a person").
- **Faith mode (hidden)**. Opened only by press-and-hold on the headword. No toggle, no label, address bar and tab title unchanged. Six blocks in fixed order: Parallel (an analogy, not the same doctrine) → Difference (no winner) → Christian bridge (two gaps only) → Linguistic root → Historical timeline → Verse list. Empty blocks say "The dictionary does not cover this yet."
- **Language rules**. Use mūrti, sacred narrative, deva, saṃsāra, rebirth, Theravāda. Never idol, myth, demigod, Hindu trinity, Hīnayāna, "heaven" for nirvana. Trimūrti ≠ Trinity; avatāra ≠ Incarnation; anattā ≠ soul. Name the school when schools disagree.
- **Human in the loop**. Experts edit drafts → "Waiting for review" → a reviewer approves; the public page never shows the expert's name.
- **Missing word** stays on the dictionary: "Not in the dictionary." No AI draft.

Layer 2 — Monthly question & Ideas map (public, anonymous)

- One question a month about one faith word (salvation, marriage, love, faith); the word links to its page. No account, not tied to a pastor. Closes at month end.
- Answers pass the moderation pipeline (§9), then are grouped into concepts.
- **Admin draft** (quiet footer link): counts only, flagged items held, duplicates/noise counted not shown. **Publish** is the only thing that updates the public map.
- **Ideas map**: a dedicated screen. Bigger = mentioned more. This-month / last-month scrub; new concepts marked. Tap a concept → label, count, neighbours. No sentences, no names; concepts under 3 hidden.

Layer 3 — Pastor care (private)

- **Pastor home**: stage ring, one next step, handwritten mentor note, Tracks · Pack · Church app.
- **Check-in (when you can)**: mood compared with a usual day, two short boxes. Result screen shows one outcome: *Steady* (one encouragement), *Hard note* ("A person will see this." + crisis line + Message your mentor), *Failed* (one sentence + Retry), *Saved offline*.
- **Reviewer**: list of code names → Preparing (4 steps, then collapses) → review with three blocks (month counts, encouragement, one held item "A person still has to decide") → Acknowledge → "Review still open."
- **Leader alert**: pauses Faith mode, monthly answers, check-ins and integrations in a region. Typed confirm. Enforced on the server.
- **Church apps (US only, optional)**: Planning Center → sign-in → Connected → imported rows (no finance) → read-only embed card for the partner app.

4. Screen inventory — what each screen does and why

Every screen passes the test "This screen lets you ___" with one primary (olive-gold) action.

Screen	Lets you...	Why it's built this way
Landing	Understand the app and start	Hero shows a real word working; tour plays inline with no player chrome; Sign in + Get started in header; footer lists guest, account, churches, privacy.
Tour video	See every module in 60–90 s	Autoplays, muted, loops, one faint pause button (WCAG 2.2.2). 16:9 and 9:16 cuts.
Onboarding 1–4	Get the idea in four calm scenes	Bridge · Conversation · Shared table · Paths meeting. No religious buildings or symbols. Back on every step; motion starts immediately; uses all three brand colours.
Welcome / Sign in / Register	Choose guest or account	Guest-first; "Continue as guest" equal in clarity; some regions guest-only.
Search	Look up a word	Field pinned to the bottom on phone (thumb reach); voice input; "Start here" words.
Word page	Read one word across traditions	Dictionary rhythm (headword → sense → sources); trust line first.
Faith mode	Go deeper, privately	Hidden for safety in sensitive places; fixed order prevents debate tone.
Save as image	Keep a word card	Private sharing without a social feed.
Monthly question	Answer once, anonymously	Permission-first wording; "Did we read you right?" correction; crisis line on crisis words.
Ideas map	See what people noticed	Own screen (not mixed with drafts); scrub + New badges; labels always readable; pinch/zoom.
Admin · Map draft	Check and publish	Two counts only; publish needs typed confirm; spike warning.
Admin · Accounts / Faith review	Manage roles; approve edits	Role changes never move a stage; names only via logged reveal.
Settings / What we store	See and delete your data	Per-item delete; delete account; theme, text size, reduce motion, language.
Pastor home	See where you are and what's next	One next step; mentor note in handwriting font; season ring.

Screen	Lets you...	Why it's built this way
Check-in + result	Tell someone how today was	No score, no risk word; edit updates the same result.
Monthly pack	See what's missing	Pieces with icons, not progress bars; "What we heard" lives here.
Tracks	Follow training	Linked from home, not previewed there (one purpose).
Reviewer queue / Preparing / Review	Decide for one pastor	Code name + broad region; AI draft labelled "Why am I seeing this?".
Leader alerts	Pause sensitive features	One switch per region; lift returns everything.
Church apps / Sign-in / Connected / Embed	Use our features in Planning Center	"How does it work?" beside Connect; demo church badge; stage never changes.
Plain page	Hide the app instantly	Close / Escape anywhere in Layer 3 or Faith mode.
Feedback board	Suggest ideas, report problems	4 kinds, merge duplicates, statuses, change log; no names.
Beta survey	Tell us how it felt (1 min)	9 easy questions, auto-advance, Back/Skip.
Analytics dashboard	See usage (admins)	Counts only, under-10 rule, alert regions hidden.
Waitlist	Join early access	One field, one role chip, "Try as guest".

Every data screen also has loading, empty (with custom illustration), error + retry, offline and unavailable states — see `#state=` in any kit URL.

5. Design system

Colour — one primary, one subtle partner, purposeful semantics

Role	Token	Use
 Ink (green-black)	<code>--ink-*</code>	Background ramp; never pure black
 Lamp (olive-gold)	<code>--lamp-*</code>	The one primary action per screen, focus, selection
 Clay	<code>--clay-*</code>	Links and human moments (mentor, thank-you)
 Bone	<code>--bone-*</code>	Text ramp
Tradition colours	equal weight	Map and labels only — none dominant
Semantic	success / warn / error / info	Only for status, never decoration; never the only signal

Why "Lamp & Clay": earthy, warm and calm — current Gen Z taste leans away from neon toward olive, sage and terracotta. It avoids the Spotify/Hulu black-and-green look and the purple-gradient "AI" look.

Type

- **Bricolage Grotesque** — display (headlines, numbers).
- **Atkinson Hyperlegible** — body, designed for low-vision readability.
- **Kalam** — handwriting, only for a mentor's note or a pastor's own words.
- **Noto Sans** Devanagari, Bengali, Myanmar, Khmer — script fallbacks.
- Six sizes only (display, H1, H2, body, small, caption); body $\geq$ 16px on phone.

Layout, spacing, elevation

- 4-pt spacing scale; equal top and side padding on every card (fixed screen by screen).
- Phone 390 first; desktop 1280 with a 12-column grid and a side rail.
- Three shadow levels; overlays use a soft linear gradient scrim, never a glow.
- Radii: 12 (inputs), 16 (cards), pill (buttons, chips).

Components & states

Button, IconButton, Badge, Tag, Card, Icon (Heroicons-style outline, sized to text), SeasonRing, PackPieces, SourceLine, form fields. **Every interactive control has 4 states** — default, hover (a shade brighter), pressed (scale .98), focus (2px lamp ring) — plus disabled. Tap targets $\geq$ 44px.

6. Motion & emotional design

Token	Value	Use
--dur-fast	120 ms	Hover, press
--dur-base	180 ms	Mode switch, reveal
--dur-slow	320 ms	Stage marker move
--ease-out	cubic-bezier(.2,.7,.2,1)	Default

- **Motion only explains a change.** No scroll fade-ins, particles, cursor glows or scroll-jacking.
- **Signature moments** (unique to us): headword press-and-hold reveal; ideas settling into the map after publish; mentor note "written" in; season ring turning on a stage change; calm onboarding scenes.
- **Emotional return loop** without pressure: one question a month, "what changed since last month", a handwritten note — no streaks, no badges, no guilt reminders.
- **Reduced motion** respected everywhere; the tour starts paused.
- The split-and-rejoin animation was removed at your request.

7. Onboarding, guest-first, landing, tour

- **Onboarding** (4 scenes): a bridge between two sides; diverse people in conversation; a shared table; several paths meeting. Avoided: crosses, churches, mosques, temples, holy books as hero, single religious clothing, "one path" metaphors.
- **Guest mode:** look up, save on device, answer monthly, view map, "Did this help?". Upgrading keeps saved words; past answers stay anonymous.
- **Landing:** one purpose (look up a word); tour inline; no repeated sections (de-duplicated after review).
- **Tour video:** covers every module incl. alert mode and Planning Center; marketing angle — Gen Z-designed, strengthens faith, best explanation — without revealing hidden layers to the public.
- **Demo vs. public:** judges see everything via the demo role switcher; public users never see hidden features.

8. Guardrails & safety (server-enforced)

Guardrail	Rule	Why
Prompt injection	"Ignore your rules", "system prompt", "act as"... → blocked, request id, nothing sent to a model	The assistant must never be steered by user text
Harmful / hateful content	Blocked or held for a person	Protect traditions and people
Personal details	Names, phones, emails, addresses removed before storage	Anonymity in sensitive regions
Crisis language	Monthly: support line. Check-in: a real mentor notified	Support, never label
Concept map	Only concepts with count ≥ 3 ; no text, names, places, times	Prevent re-identification
Tone check	Ranking language withheld until a person fixes it	"No winner" promise
Coverage	No verified source → "not covered yet", nothing generated	No hallucinated doctrine
Model text	Rendered as plain text, never HTML	Security
Why this?	Every assistant-written text has a "Why am I seeing this?" line	Transparency (Microsoft HAX G11)
Stage	Only a person moves a pastor's stage; logged	Human decision

9. Moderation pipeline (every typed input)

1. **Limits** — answer 3–280, check-in ≤ 500 , edits 10–1200 chars.
2. **Unclear** — emoji only, "aaaaa", keyboard runs $\rightarrow$ "That doesn't read as words yet."
3. **Personal details** removed.
4. **Sensitive categories** — self-harm (person within 24 h), hate (hold), sexual / threat (block).
5. **Duplicates** — one answer per device per month; changes replace; bursts from one network counted once.
6. **Rate** — 5 sends per device per hour.
7. **Group** into ideas; "Did we read you right?"
8. **Hide small** — ideas < 3 never publish; counts < 10 read "under 10".
9. **Spike** — idea rising $> 3\times$ from one region in an hour flagged.
10. **Person review** — rename, merge, remove, publish.
11. **Noise** — small Laplace noise ($\epsilon \approx 1$) on published counts: "Counts shift slightly to protect people."

How do we know it's noise? Rules (length, repeats, no vowels, keyboard runs) catch most junk on the device; the server adds a model check and a person reviews anything held. Noise is counted, never shown.

10. Privacy, alert mode, hidden features

- **What we store** page with per-item delete; delete account leaves answers only as counts.
- **Quick-hide:** Close / Escape → a plain page, from any Layer 3 screen and Faith mode.
- **Alert mode:** a leader raises an alert for a region (typed confirm). Paused: Faith mode, monthly answers, map publish, check-ins, integrations, voice, save-image. Calm message, no alarm. Lift → all return. Server returns `503 feature_off` ; the UI hiding is not enough.
- **Hidden vs. open:** open in the normal flow — dictionary, monthly, map, guest. Hidden in sensitive areas — Faith mode, pastor layer. Shown in full to judges in the demo.
- Low profile by design: no public sharing, no video, no voice recording stored, no social feed.

11. Languages & low bandwidth

- English, हिन्दी, বাংলা, नेपाली, ଓଡ଼ିଆ, ୟିଡ଼ୁଂ across all three layers; tradition terms stay as written. Native-speaker review required before launch.
- Language picker restyled (dark list, no browser-default white).
- Plain words for young readers: technical words replaced (e.g. "lexicon" → "dictionary", "packet" → "review", "integrations" → "church apps").
- Offline: check-ins queue on the phone and send later, no duplicates.
- Demo mode: no spinners in front of judges.

12. References & how we used them

Rule: **borrow the idea, never the look.** Nothing changed our palette, type or icons.

Screen	Reference	What we borrowed
Word page	Merriam-Webster; Perplexity	Dictionary rhythm; trust line before text
Faith mode / verses	Substack, Matter (dark readers)	Label → title → meta → body reading rhythm
Public voice	Gen Z Bible	Dark, bold, phone-first; no mood menus or chat
Search (phone)	Moonlitt	Field pinned for the thumb
Monthly question	Oura; Tiimo	Permission-first question; previous-questions list
Ideas map	Stardust; Pol.is; Neo4j Bloom	"What changed" sentence; shared vs. apart; pan/zoom canvas
Check-in	Oura; Learn2Care; Tolan	Rows vs. usual day; result states; one warm sentence
Reviewer	Canny Autopilot; Linear	Source → suggestion side by side; one stage path
Church apps	Supabase; Planning Center	"How does it work?" beside Connect; calm admin density
Onboarding	Mobbin web onboarding flows	Short steps, Back on every step, one idea per step
Icons	Heroicons	Outline style, sized to text

13. Rules we followed

"Not AI slop" (video transcript + our checklist)

Do

- Show a real screen working
- H1, one subtext line, one action
- One primary + one secondary button
- Text 4.5:1 on its surface
- Hover one shade brighter
- Motion only to explain a change
- Every state drawn
- Name the visual problem, then fix it

Don't

- Icon + title cards, bento grids
- Fake dashboards
- Eyebrow label on every card
- Purple gradients, emoji icons
- Random glows / lights
- Uneven padding, lazy border-only selected states
- Fade-in on scroll, particles
- Inter / Roboto / overused fonts

Test: would we have spent a week hand-building this? If not, cut it.

UI/UX fundamentals checklist (your 12-point brief)

Signifiers · visual hierarchy (most important at top) · icons where meaningful · grid & layout · standard structure · 6 type sizes · one primary colour + subtle partner (colour ramp) · semantic colours for purpose · shadows · icons/buttons sized to text · 4 button states · micro-interactions · linear-gradient overlays.

One purpose per screen

A screen may hold many things that serve one job. Split when tasks are unrelated or sequential. Applied to monthly, map, pastor home, hard-note result, onboarding step 4 and the map/draft separation.

YC / startup design guidance

Ship the simplest thing that works; clarity over cleverness; design for the user's first 10 seconds; remove before adding.

Microsoft Human-AI guidelines (18) & Google People + AI

Make clear what it can do (tour, onboarding) · how well ("It can be wrong") · social norms and bias (banned words, own texts) · easy to dismiss (Escape) · easy to correct (Did we read you right?, expert edits) · scope down when unsure (not covered) · explain why (sources, Why this?) · remember recent (saved words) · give control (delete, reduce motion).

Baymard principles (forms, search, checkout-style flows)

Inline validation after leaving a field; errors say how to fix; keep input on error; minimal fields; clear field labels (not placeholder-only); search tolerates spelling and diacritics; no forced account.

Accessibility (WCAG 2.2 AA)

Contrast 4.5:1 / 3:1; tap $\geq$ 44px; names on icon buttons; heading order; focus rings; keyboard-only paths; reduced motion; pause for moving content; colour never the only signal; text up to 200%.

14. Testing we did

Results are generated live each time a report is opened, so this table lists what each test checks and where to run it. Copy your latest scores in before presenting.

Order: X-ray & audits → automated tests → AI-assisted review → traditional heuristic testing → strict benchmark → human tests (pending).

Test	What it checks	Result
UI X-ray (mobile 390, desktop 1280)	Outlines every element: overlaps, overflow, sideways scroll, clipped text, alignment	See report
Accessibility audit	Contrast, tap < 44, unnamed controls, headings, tiny text	See report
Phone / keyboard / contrast	No sideways scroll; focusable; contrast	See report
Component & input rules (20 cases)	Limits, unclear text, personal details, crisis, duplicates, rate, injection	See report
Every screen loads (smoke)	Landing, onboarding, app, pastor, reviewer, tour	See report
Nielsen's 10 heuristics	Scored 1–5 each; pass = avg ≥ 4, none < 3	See report
Cognitive walkthrough (6 tasks)	Know what to do / see it / connect to goal / know it worked	See report
Consistency audit	Same thing, same name, same shape	See report
Error & edge states	Empty/bad input announced, fixable, input kept	See report
Navigation & structure	Way back, clear title, no dead end	See report
Microcopy	Flesch–Kincaid ≤ 8, long sentences, vague buttons	See report
Keyboard only	Tab reach, no positive tabindex, visible focus	See report
Reduced motion	No looping animation after 2.5 s	See report
AI-assisted review	One purpose, plain language, banned words, fairness, no sales tone	See report
Messy-data stress	Long names, long regions, many items	See report
Silent-break guard	Fingerprint of every screen vs. baseline	See report

Why so many layers? AI reviewers tend to say "everything passes". We added a strict, no-context reviewer and a non-AI fingerprint guard so nothing breaks silently — and kept human tests as the final word.

15. Benchmarks & scores

- **Strict vs. informed heuristic score:** strict (harsh, no context) is the number we report; a big gap flags screens to check by hand.
- **Benchmark vs. YouVersion:** same strict reviewer and format run on YouVersion's public web pages (bible.com), 3 runs averaged. Scores are produced live in `ui_kits/Benchmark vs YouVersion.html` ; copy the numbers from your run.
- **NN/g 7-step product benchmark:** steps 1-2 done (7 top tasks, metrics and targets); steps 3-7 need people.

Metric	Target
Task success (unaided)	$\geq 90\%$ per task
Time on task	$T1 \leq 30\text{ s} \cdot T2 \leq 60\text{ s} \cdot T5 \leq 90\text{ s}$
SEQ	$\geq 5.5 / 7$
SUS	≥ 75 (68 average)
Fairness	$\geq 4 / 5$ in every tradition
Felt pushed to convert	$\leq 10\%$

16. Originality check

Each screen was compared with the closest top-app screens on Mobbin. **0 high risk.** Two to watch: dark + green family (Spotify) — answered by green-black ink, olive-gold, clay; Ideas map vs. Obsidian graph — answered by monthly published picture, size = mentions, scrub, equal tradition colours, no physics or glow. Six signature moments nothing else has (press-and-hold reveal, map scrub, handwritten mentor note, "A person will see this", region alert, guest-first faith dictionary).

17. Beta, feedback, analytics

- **Feedback board:** Idea · Missing word · Something is wrong · Not fair to a tradition; 8–280 chars; blocks names/numbers/places; crisis → support line; duplicates merge; 3 a day; statuses Being read → Planned → Building → Shipped; public change log.
 - **Beta survey** (standalone, 1 min, 9 questions): role · ease 1–5 · helped understand · fair · pushy · 3 feel-words · use again · broken (optional) · follow-up email (optional). Config `endpoint` for Cursor.
 - **Analytics dashboard** (admins): visitors, words looked up, monthly answers, 7-day return; daily visits; first-visit funnel; top words; languages; monthly moderation counts; check-in outcomes; feedback; survey scores (SUS-scale, NPS, fairness, pushed). Counts only, under-10 rule, alert regions hidden, Faith mode never tracked.
 - **Human test plan:** 46-question survey (screening, 5-second, first-click, tasks + SEQ, SUS, fairness, feeling, privacy, role-specific).
-

18. Backend contract & handoff

The `contract/` folder is the single source of truth for Cursor: `api.md` (every endpoint), `types.ts` (shapes), `guardrails.md`, `moderation.md`, `demo-flow.md` (34-step flow and demo data), `launch.md` (waitlist, analytics, email, store, launch). Rule: if a field isn't in `types.ts`, it doesn't leave the server.

Live demo: `Live Demo.html` — 10 steps with "Say this" lines, phone frame, role switcher (Guest, Admin, two pastors, Reviewer, Leader, US church), failure switches, Reset. Fictional place names. Deploy with `vercel.json` (see `DEPLOY.md`).

19. Human flourishing by design

We used the six domains of human flourishing (Harvard Human Flourishing Program: happiness & life satisfaction, health, meaning & purpose, character & virtue, close social relationships, financial & material stability) as a filter. A feature stays only if it supports one of them **without** taking attention, money or dignity.

Domain	How the design supports it	What we refused to build
Meaning & purpose	One word at a time; "shared human questions"; the monthly question invites reflection, not performance	Infinite feed, trending lists
Close relationships	Mentor note in handwriting; "A person will see this"; Message your mentor	Chatbot companion standing in for people
Health (mental)	Check-ins "when you can"; rest-focused encouragement; crisis line; reduced motion	Streaks, daily quotas, red badges, guilt reminders
Character & virtue	Fair description of every tradition; no ranking; "Not fair to a tradition" feedback kind	Debate scores, "who's right" framing
Happiness & satisfaction	Calm palette, slow motion, finishing moments that feel complete	Gamified points, confetti, dark patterns
Stability & safety	Guest-first, quick-hide, alert mode, no finance collected, works offline	Data selling, public profiles, location sharing

Clean and easy because of it

- **One purpose per screen** — less to decide, less stress.
- **Fewer choices, better defaults** — guest by default, language remembered, auto-advance in surveys.
- **Agency** — Back on every step, Skip on optional, delete anything, undo where it matters.
- **Dignity** — no scores on people, code names, "Why am I seeing this?".
- **Completion, not compulsion** — sessions end cleanly; nothing pulls you back except next month's question.
- **Time well spent** — success is "understood a word", not minutes in app.

Measure it. Beta survey asks "Did it help you understand a word?" and "Did it feel pushy?"; dashboard tracks return next month, not daily active minutes.

20. Open build & license

- This document, the design system, screens and contract are published under a permissive license: **code MIT, docs and design CC BY 4.0** (see `LICENSE.md`).
- Builders can start from our **work** — tokens, components, screens, moderation rules, test suites — not just our conclusions.
- Excluded from the license: tradition content needing source permissions (verse texts are licensed separately), and user data (none is shipped).
- Publish: repo + Vercel demo link + this PDF.

21. Shared schema & MCP server

The capability other teams can call is the **fair cross-tradition word lookup** and the **safe anonymous input pipeline**. Both are defined once in `contract/types.ts` and exposed three ways:

Surface	What
REST	<code>GET /terms?q=</code> , <code>GET /terms/:term</code> , <code>POST /moderate</code> , <code>GET /map/:month</code> (published counts only)
MCP server	Tools in <code>contract/mcp-tools.json</code> : <code>lookup_word</code> , <code>compare_traditions</code> , <code>moderate_text</code> , <code>get_ideas_map</code> . Read-only; no user data; alert regions return <code>feature_off</code> .
Embed	Read-only card for church apps (Planning Center demo)

Faith mode, pastor data and reviewer tools are never exposed externally.

Status. The schema is published. The endpoint and MCP server are specified but not deployed yet — say this plainly in the presentation.

22. Published eval set & rubric

Other teams can run their own agents and UIs against our checks. Files in `evals/` :

- **rubric.md** — fairness & tone rubric for any interfaith explainer (6 criteria, 1–5, pass ≥ 4 , no criterion < 3) + strict UI heuristic rubric (Nielsen 10, harsh/no-context scoring).
- **cases.json** — 30 input cases: limits, gibberish, personal details, crisis, hate, injection, duplicates, banned terms ("idol", "Hīnayāna"), conflation traps (Trimūrti $\neq$ Trinity), ranking language — each with expected outcome.

How to use. Feed each case to your agent; compare outcome to `expect` . Report pass rate and the strict UI score from 3 runs averaged.

23. Reusable patterns (named & documented)

Pattern	What it is	Transfers to
Side-by-Side, Never Ranked	One concept described by each tradition from its own sources, equal weight, no winner	Politics, medicine, law, any contested topic
Hold to Reveal	Sensitive layer behind press-and-hold, no visible sign, clean return	Safety apps, journalism, abuse support
Publish-Gated Picture	Public aggregate updates only when a person publishes; small groups hidden; noise added	Civic surveys, HR pulse, research
A Person Will See This	Hard input → one honest line + a human route, no score or risk label	Health, education, employee wellbeing
Region Alert	One switch pauses sensitive features per region, enforced server-side	Any app operating under censorship or conflict
Why Am I Seeing This	Every AI-written text carries a plain one-line provenance	Any AI product
Strict + Fingerprint Testing	Harsh no-context AI review + non-AI screen fingerprint to stop silent breaks	Any AI-assisted design team

24. Practitioner evidence

To fill during the event. Only real, attributable quotes — no invented ones.

Who (role, with permission)	What they did	What they said (exact)	What we changed

Capture with the beta survey (`?from=event`) and a 2-line consent note.

25. What we tried that didn't work

Tried	Why we dropped it
Normal / Faith toggle on the word page	Visible label put people at risk; replaced by hidden press-and-hold
Comments and votes for guests	Invites debate and ranking of traditions; against "never ranked"
Daily check-ins with counts ("22 of 30")	Felt like surveillance to pastors under pressure; now "when you can"
Finance totals from Planning Center	Sensitive, little pastoral value
Split-and-rejoin animation	Read as "division"; removed
Map and draft on one screen	Overlaps, messy; split into a dedicated map screen
Light moss-green labels, then white	Hard to read / harsh; settled on black on green
Landing with repeated sections + "click to watch" tour	Repetition; tour now plays inline
Video player chrome on the tour	Looked like a player, not a product; one faint pause button
AI-only testing	Said everything passed; added strict reviewer + fingerprint guard
Auto-generated PNG exports	Renderer missed icons/iframes; use Chrome full-size screenshots
Single-file offline bundle of the demo	Nested screens didn't load; deploy the folder instead
Lazy-loaded screens on the All Screens board	Stayed blank on the zoom canvas; now load in a queue
Browser-default language dropdown	White list on dark UI; custom dark picker
Video sharing / public voice posts	Raises the app's public profile in sensitive regions; kept voice input only, nothing stored

26. What's left

- Human tests (task success, SEQ, SUS, fairness by tradition) — this week.
- Native-speaker review of all 5 translations.
- Backend: moderation, alert enforcement, publish, noise, real analytics and survey storage.
- Connected church rows from the real Planning Center API.
- PIN lock for shared phones; long-text and 200% text pass on real devices.

27. File map

File / folder	What it is
Live Demo.html · Full Demo.html · All Screens.html	Judge demo · every screen board · phone + desktop canvas
ui_kits/landing · tour · public · pipeline	Landing, tour video, Layers 1–2, Layer 3
ui_kits/feedback · survey · analytics · waitlist	Beta tooling
ui_kits/xray · audit · a11y · tests · Expert Review · More Tests · Benchmark*	Test suites
UX Benchmark · UX Testing Research · Similarity Check · One Purpose · Mobbin Changes	Research and decisions
tokens/ · styles.css · components/ · guidelines/	Design system
contract/ · DEPLOY.md · vercel.json	Backend handoff and deployment
LICENSE.md · evals/ · contract/mcp-tools.json	Open license, eval set, MCP tool schema
exports/beta-survey	Standalone survey for Cursor